

Project Alpha - internal notes

This sample PDF exists so you can run:

```
python src/ingest.py
```

The RAG pipeline extracts text from PDFs using pypdf. Use real PDFs with embedded text (exported from Word, Google Docs, or printed from Markdown). Scanned pages need OCR first; image-only PDFs usually return no text.

Key facts for testing retrieval:

- Project code name: Alpha
- Owner team: Platform
- SLA target: 99.9% monthly uptime